

MASR — Meta Abstract Semantics Representation

Brian Beckman

23 Apr 2023

Contents

1	Abstract	3
2	Summary and Cheat Sheet	4
2.1	Specs	4
2.2	Namespace-Qualified Keywords	4
2.3	Three Kinds of Specs	4
2.4	Full-Form	4
2.5	Sugar	4
2.6	Terms and Heads	5
2.7	Multi-Specs	7
2.8	Nested Multi-Specs	8
2.9	Light Sugar, Heavy Sugar	9
3	Issues with ASDL	13
3.1	ASDL is not Type-Checked	13
3.2	ASDL is Moribund	13
3.3	ASDL is Incomplete	13
3.4	ASDL's ASR is Volatile	13
3.5	ASDL is Ambiguous	13
4	Clojure Solves ASDL Issues	14
4.1	Clojure Is a Complete Fit for ASR	14
4.2	Clojure is not Moribund	14
4.3	Clojure Has Type-Checking Tools	14
4.4	ASR in Clojure Solves Volatility	14
5	MASR Definitions	15
6	MASR Tenets	17
7	Base Specs	18
7.1	Atoms: <code>int</code> , <code>float</code> , <code>bool</code> , <code>nat</code>	18
7.2	Notes	19

8	Term-Like Nodes	20
8.1	dimensions [<i>plural</i>]	20
8.2	identifier [<i>singular</i>]	22
8.3	identifiers [<i>plural</i>]	24
9	Specs	27
9.1	intent	28
9.2	TODO unit	35
9.3	TODO symbol	36
9.4	storage_type	38
9.5	access	39
9.6	TODO deftype	39
9.7	presence	39
9.8	abi	40
9.9	TODO stmt	44
9.10	TODO expr	45
9.11	ttype	46
9.12	TODO restriction_arg	46
9.13	TODO binop	47
9.14	TODO logicalbinop	48
9.15	TODO cmpop	49
9.16	TODO integerboz	50
9.17	TODO arraybound	51
9.18	TODO arraystorage	52
9.19	TODO cast_kind	53
9.20	dimension	54
9.21	TODO alloc_arg	56
9.22	TODO attribute	57
9.23	TODO attribute_arg	58
9.24	TODO call_arg	59
9.25	TODO tbind	60
9.26	TODO array_index	61
9.27	TODO do_loop_head	62
9.28	TODO case_stmt	63
9.29	TODO type_stmt	64
9.30	TODO enumtype	65
10	Implicit Terms	66
10.1	symtab_id	66
10.2	TODO symbol_table	67

1 Abstract

Abstract Semantics Representation (ASR) is a novel intermediate representation (IR)¹ for a new collection of LCompilers [sic].² ASR is agnostic to the particular programming language under compilation. Current compiler front-ends targeting ASR include LFortran³ and LPython.⁴ ASR is also agnostic to the back end. ASR currently targets LLVM, x86, C, and WASM⁵

Typical IRs encode semantics as decorations on the Abstract Syntax Tree, (AST)⁶ ASR lifts *semantics* to the top level and expunges the syntax of the surface language as early as possible. Free of syntactical baggage, ASR optimizers are cleaner and faster than average, and ASR back ends are completely reusable. If syntax information is ever necessary, as with semantical-feedback parsing, such information will be encoded as decorations on the ASR, rather than the other way around.

Current specifications for ASR are written in ASDL,⁷ a metalanguage similar in spirit to yacc but less rich, by design.⁸ To build an LCompiler like LFortran or LPython, the ASDL grammar⁹ for ASR is parsed, and a library in C++, libasr,¹⁰ is generated. Compiler front ends call functions in this library to manipulate ASR and to emit code from the back ends.

ASDL has several deficiencies, and MASR¹¹ alleviates them. Chief among the deficiencies is the lack of type-checking. MASR adds a type system to ASR via Clojure *specs*.¹² MASR is a complete programming language in its own regard. It is, in fact, a Domain-Specific Language (DSL),¹³ embedded in Clojure.¹⁴

We aim to replace ASDL with MASR and to integrate MASR with the LCompiler code base. When so integrated in the future, MASR will be called LASR.

This document is pedagogical, both explaining MASR and teaching how to extend and maintain its Clojure code.

This document may lag the Clojure code. It may also lag libasr, at least until MASR replaces ASDL. The document mirrors an ASDL snapshot.⁹

¹https://en.wikipedia.org/wiki/Intermediate_representation

²<https://github.com/lcompilers/libasr>

³<https://lfortran.org/>

⁴<https://lpython.org/>

⁵<https://webassembly.org/>

⁶https://en.wikipedia.org/wiki/Abstract_syntax_tree

⁷https://en.wikipedia.org/wiki/Abstract-Type_and_Scheme-Definition_Language

⁸<https://en.wikipedia.org/wiki/Yacc>

⁹https://github.com/rebcabin/masr/blob/main/ASR_2023_APR_06_snapshot.asdl

¹⁰<https://github.com/lfortran/lfortran/tree/c648a8d824242b676512a038bf2257f3b28dad3b/src/libasr>

¹¹pronounced “maser;” it is a Physics pun

¹²<https://clojure.org/guides/spec>

¹³https://en.wikipedia.org/wiki/Domain-specific_language

¹⁴<https://en.wikipedia.org/wiki/Clojure>

2 Summary and Cheat Sheet

2.1 Specs

Clojure specs double as *types*, constituting ASR's type system.¹² The functions `s/valid?` and `s/conform` recursively check instances of a form against specs.

2.2 Namespace-Qualified Keywords

Specs are named, defined, and registered via *namespace-qualified keywords*. All MASR spec keywords are registered in namespace `masr.specs`. The file `specs.clj` defines the namespace `masr.specs`. In the file `specs.clj`, a double-colon shorthand is available. For example, `::nat` in the file `specs.clj` is short for `:masr.specs/nat`.

2.3 Three Kinds of Specs

MASR has three kinds of registered specs:

simple specs — registered via `s/def`, as in `(s/def ::bool boolean?)`

entity specs — registered via `s/keys`; have required and optional attributes; examples below

multi-specs — defined via `s/multi-spec`, registered via `s/def`; multi-specs have a distinguished *tag* attribute like tagged unions in C; examples below

MASR multi-specs are tagged collections of entity specs.

2.4 Full-Form

Full-form instances that are checked against specs are Clojure *hash-maps*:¹⁵ collections of key-value pairs like Python dictionaries. For example,

```
;; key          value
{:::term        ::intent,
 ::intent-enum 'Local}
```

In MASR, all keys in all hash-maps are namespace-qualified keywords. Such keys may have specs registered for them, or not. When a spec is registered for a key, automatic recursive type-checking is invoked.

2.5 Sugar

Every spec *qua* type has a full form as well as several shorter sugared forms. Sugar is defined by functions like `Integer` and `Integer-` that return instances in full-form. Sugar comes in two flavors, *light* and *heavy*. See Section 2.9.

¹⁵<https://clojuredocs.org/clojure.core/hash-map>

2.6 Terms and Heads

MASR defines *terms* and *heads* that describe the semantics of programs. Terms are top-most in the ASR grammar and heads are at the bottom level. There are only two levels.

The following tables summarize this document via conforming examples, written in the recommended sugar form.

Equally important are non-conforming examples. See the body of this document, tests in `specs.clj`, and `deftest` in `core_tests` for many non-conforming examples.

Table 1: Atomic and Naked Specs: No Sugar

Spec	Predicate	Link	Example
::bool	boolean?	7.1	true
::float	float?	7.1	3.142
::int	int?	7.1	-1789

Table 2: Top-Level *term-like* Specs, not in ASDL

Spec	Link	Example
::nat	7.1.1	(nat 42)
::identifier	8.2	(identifier 'boofar)
::identifier-set	8.3	(identifier-set ['a 'a])
::identifier-list	8.3	(identifier-list ['a 'a])
::identifier-suit	8.3	(identifier-suit ['a 'b])
::dimensions	8.1	(dimensions [[6 60] [1 42]])

Table 3: Higher-Order Specs, Defined in Terms of Other Specs

Spec	Link or Spec	Example
::syntab-id	10.1	(syntab-id 42)
::value-attr	::bool	(value-attr false)
::dependencies	::identifier-set	(dependencies ['a 'b 'c])
::symbolic-value	TODO empty?	(symbolic-value ())
::value	TODO empty?	(value ())
::type-declaration	TODO nilable syntab-id	(type-declaration nil)
::varnym	::identifier	(varnym 'x)

Table 4: *Term* Multi-Specs; Like Tagged Unions

Multi-Spec	Term	Link	Example
::asr-term	::dimension	9.20	(dimension [6 60])
::asr-term	::intent	9.1	(intent 'Local)
::asr-term	::storage-type	9.4	(storage-type 'Default)
::asr-term	::abi	9.8	(abi 'Source)
::asr-term	::access	9.5	(access 'Public)
::asr-term	::presence	9.7	(presence 'Required)

Table 5: ::asr-term Specs with Nested *Head* Multi-Specs

Term	Head	Link	Example
::ttype	::Integer	9.11	(Integer 4 [[6 60] [1 42]])
::ttype	::Real	9.11	(Real 8 [[6 60] [1 42]])
::ttype	::Complex	9.11	(Complex 4 [[6 60] [1 42]])
::ttype	::Logical	9.11	(Logical 1 [[6 60] [1 42]])
::symbol	::Variable	9.3.1	(Variable 42 'x (Integer 4)...))

2.7 Multi-Specs

Instance hash-maps that conform to multi-spec `::asr-term` are polymorphic. They have a tag attribute, fetched via `::term`, that must match a term defmethod.¹⁶ For instance, the following example is a valid `::asr-term` in full-form; its `::term` attribute is `::intent`:

```
(s/valid? ::asr-term
  {::term      ::intent, ;; matches a defmethod
   ::intent-enum 'Local}) ;; specifies contents
```

where `::intent-enum` is a simple spec defined and registered via `s/def`:

```
(s/def ::intent-enum ;; #{...} is a Clojure set.
  #{'Local 'In 'Out 'InOut 'ReturnVar 'Unspecified})
```

Its `::term`, `::intent`, matches a term defmethod below. The keyword, `::term`, doubles as an attribute key in the instance and as a function for fetching the `::term` value from an instance.

Here is another `::asr-term` in full-form, matching a term defmethod for `::abi`:

```
(s/valid? ::asr-term
  {::term      ::abi, ;; matches a defmethod
   ::abi-enum   'Source ;; specifies contents
   ::abi-external false}) ;; specifies contents
```

where

```
(s/def ::abi-external ::bool)
```

Other `::asr-term` specs follow the obvious pattern. The `::term` attributes, `::intent`, `::abi`, etc., each match a term defmethod:

```
(defmulti term ::term) ;; ::term fetches the tag-value
(defmethod term ::intent [_] ;; tag-value match
  (s/keys :req [::term ::intent-enum])) ;; entity spec
(defmethod term ::dimension [_] ,,,) ;; tag-value match
(defmethod term ::abi [_] ,,,) ;; tag-value match
(defmethod term ::ttype [_] ,,,) ;; tag-value match
(defmethod term ::symbol [_] ,,,) ;; tag-value match
;; etc.
```

Finally, the multi-spec itself is named `::asr-term`:

```
;;      name of the mult-spec    defmulti  tag fn
;;      -----
(s/def ::asr-term (s/multi-spec term ::term))
```

¹⁶<https://clojuredocs.org/clojure.core/defmulti>

2.8 Nested Multi-Specs

Contents of multi-specs can, themselves, be multi-specs. Such cases obtain when an `::asr-term` has multiple function-like heads. Examples include `::ttype`, `::symbol`, `::expr`, and `::stmt`.

The names of all multi-specs in MASR, nested or not, begin with `::asr-` and end with either `term` or `<some-term>-head`. Examples: `::asr-term` and `::asr-ttype-head`. There is only one level of nesting: terms above heads.

Here is the `::asr-term` entity spec for `::ttype`. The *nested* multi-spec is named `::asr-type-head`.

```
(defmethod term ::ttype [_]
  (s/keys :req [::term ::asr-ttype-head])) ;; entity spec
```

where

```
(defmulti ttype-head ::ttype-head) ;; tag fetcher
(defmethod ttype-head ::Integer ,,,) ;; tag match
(defmethod ttype-head ::Real ,,,) ,,,
(s/def ::asr-ttype-head ;; name of the multi-spec
  ;; ties together a defmulti and a tag fetcher
  ;;
  ;;
  (s/multi-spec ttype-head ::ttype-head))
```

Here is a conforming example in full-form:

```
(s/valid? ::asr-term
  {::term ::ttype,
   ::asr-ttype-head
   {::ttype-head ::Real, ::real-kind 4,
    ::dimensions [[6 60] [1 42]]})
```

Likewise, here is the `::asr-term` spec for `::symbol`:

```
(defmulti symbol-head ::symbol-head)
(defmethod symbol-head ::Variable [_]
  (s/keys :req [::symbol-head ::symtab-id ::varnym ,,,]))
(defmethod symbol-head ::Module [_] ,,,)
(defmethod symbol-head ::Function [_] ,,,) ,,,
(s/def ::asr-symbol-head
  (s/multi-spec symbol-head ::symbolhead))
```


Here is a conforming example for `::Variable` in full-form, abbreviated:

```
(s/valid?
  ::asr-term {::term ::symbol,
    ::asr-symbol-head
    {::symbol-head ::Variable
      ::symtab-id (nat 2)
      ::varnym (identifier 'x)
      ::intent (intent 'Local)
      ::ttype (ttype (Integer 4 [[0 42]]) ,,, )})
```

2.9 Light Sugar, Heavy Sugar

Light-sugar forms are shorter than full-form, but longer and more explicit than *heavy-sugar*. Heavy-sugar require positional arguments, and are thus more brittle, especially for long specifications with many arguments, where the risk is high of getting arguments out of order in hand-written code.

Light-sugar specs are returned by functions like `Integer-` whose names have trailing hyphens. Light sugar functions typically have keyword arguments, partitioned into required and optional-with-defaults. The keyword argument lists of light-sugar functions do not depend on order. The following two examples conform to both `::asr-term` and to `::ttype`:

```
(ttype (Integer- {::dimensions [], :kind 4}))
(ttype (Integer- {::kind 4, :dimensions []}))
```

Heavy-sugar specs are returned by functions like `Integer` whose names do not have trailing hyphens. Heavy-sugar specs are compatible with current `libasr --show-asr` syntax. Heavy-sugar functions employ positional arguments that depend on order. Final arguments may have defaults. For example, the following examples conform to both `::asr-term` and to `::ttype`:

```
(ttype (Integer))
(ttype (Integer 4))
(ttype (Integer 2 []))
(ttype (Integer 8 [[6 60] [1 42]]))
```

Here is a conforming spec for `::Variable` in heavy sugar; it also conforms to `::asr-term`:

```
(let [a-valid (Variable 2 'x (Integer 4)
  nil [] 'Local
  [] [] 'Default
  'Source 'Public 'Required
  false)]
  (s/valid? ::asr-term a-valid) := true
  (s/valid? ::Variable a-valid) := true)
```

2.9.1 Term Entity-Key Specs

For recursive type checking, as in `::Variable`, it is not convenient for terms to conform *only* to `::asr-term`. Therefore, we define redundant *term-entity-key* specs, like `::tterm`.

Entity-key specs for asr-terms are defined as follows:

```
(s/def ::ttype
  (s/and ::asr-term ;; must conform to ::asr-term
    ;; and have tag ::ttype
    #(= ::ttype (::term %)))) ;; lambda shorthand
```

Because we have several such definitions, we write a helper function and a macro:

```
(defn term-selector-spec [kwd]
  (s/and ::asr-term
    #(= kwd (::term %)))) ;; lambda shorthand
(defmacro def-term-entity-key [term]
  (let [ns "masr.specs"
        tkw (keyword ns (str term))]
    `(s/def ~tkw ;; like ::tterm
      (term-selector-spec ~tkw))))
```

Here are some invocations of that macro:

```
(def-term-entity-key dimension)
(def-term-entity-key abi)
(def-term-entity-key ttype)
```

Term-entity-key specs establish a partition on the set of asr-terms. Every *dimension* is an asr-term. Every *abi* is an asr-term. Every *ttype* is an asr-term, and so on.

Here are some examples of extra conformance tests for sugared specs via term entity-key specs:

```
(s/valid? ::asr-term (dimension []))      := true
(s/valid? ::asr-term (dimension '(1 60))) := true
(s/valid? ::asr-term (dimension '()))     := true

(s/valid? ::dimension (dimension []))     := true
(s/valid? ::dimension (dimension '(1 60))) := true
(s/valid? ::dimension (dimension '()))     := true

(s/valid? ::asr-term (ttype (Integer 4))) := true
(s/valid? ::asr-term (ttype (Integer 4 []))) := true

(s/valid? ::ttype (ttype (Integer 4)))    := true
(s/valid? ::ttype (ttype (Integer 4 []))) := true
```

2.9.2 Enum-Like Specs

Entity-key specs are defined automatically for all *enum-like* terms via the `enum-like` macro:

```
(defmacro enum-like [term, heads]
  (let [ns "masr.specs"
        trm (keyword ns "term")      ;; like ::term
        art (keyword ns "asr-term")  ;; like ::asr-term
        tkw (keyword ns (str term))  ;; like ::intent
        ,,,]
    `(do ,,,      ;; the entity-key spec
      (s/def ~tkw  ;; like ::intent
        (s/and ~art ;; like ::asr-term
          (term-selector-spec ~tkw)))
      ,,, )))
```

Here are some examples of extra conformance tests for automatically defined term entity-keys for enum-like specs:

```
(s/valid?  ::intent (intent 'Local)) := true
(let [iex (intent 'Local)]
  (s/conform ::asr-term iex)          := iex
  (s/conform ::intent iex)             := iex)
```

2.9.3 Term-Head Entity-Key Specs

For terms like `::symbol` and `::stmt` with multiple heads like `::Variable` and `::Assignment` it is convenient to define redundant entity-key specs like the following:

```
(s/def ::Variable                ;; head entity key
  (s/and ::asr-term              ;; top multi-spec
    # (= ::Variable              ;; nested tag
      (-> % ::asr-symbol-head    ;; nested multi-spec
          ::symbol-head)))      ;; tag fetcher
  (s/def ::Assignment            ;; head entity key
    (s/and ::asr-term            ;; top multi-spec
      # (= ::Assignment          ;; nested tag
        (-> % ::asr-stmt-head    ;; nested multi-spec
            ::stmt-head)         ;; tag fetcher
```

We define these with another macro, `def-term-head--entity-key`. The definition of this macro is found in the file `specs.clj`. An example of conformance to `::Variable` is found above, in Section 2.9.

Because heads are unique in ASR, Term-head-entity-key specs establish a partition on the set of asr-terms, just like term-entity-key specs. See the section 9.3.1 on `Variable` for a test of the `::Variable` term-head-entity-key spec.

3 Issues with ASDL

Clojure solves the following issues with ASDL:

3.1 ASDL is not Type-Checked

Type-checking for ASR instances written in ASDL is expressed only in hand-written C++ code. This situation is brittle. It's much better to have a specification language for ASR that expresses a type system.

3.2 ASDL is Moribund

ASDL has not progressed since originally published in 1987. We know of no other projects adopting ASDL. We should replace ASDL with a modern metalanguage that has a robust, lively ecosystem.

3.3 ASDL is Incomplete

Much of the semantics of ASR in ASDL, beyond type-checking, is expressed only in hand-written C++ code. The reason is that ASDL is not sufficiently expressive to cover the needed cases.

As usual with such a design, it's more time-consuming and error-prone than necessary to prototype, verify, validate, visualize, modify, and debug. Something more expressive than ASDL is needed to take some responsibility off of hand-written C++ code.

3.4 ASDL's ASR is Volatile

The ASDL for ASR changes frequently, for good reasons. However, stand-aside tools like `asr-tester`¹⁷ must chase the specification. Just keeping up with ASR-in-ASDL consumes almost all development time for `asr-tester`. We should unify the language that expresses ASR with the tools that verify and test ASR.

3.5 ASDL is Ambiguous

There are many syntactic and semantic ambiguities in the ASDL grammar.⁹ For example, the type notation `integer*` might mean, in one place, a list of `integer` with duplicate entries allowed, and, in another place, a set of `integer` with duplicate entries not allowed..

ASDL is not sufficient to express such distinctions. In practice, the hand-written C++ implementations implicitly make these distinctions, hiding them from view and making them difficult to revise. It is bad practice to hide fine distinctions that have observable effects in implementations. Instead, we should express those distinctions directly in the specifications.

Because ASDL cannot express such distinctions, we must adopt something more expressive than ASDL.

¹⁷<https://github.com/rebcabin/asr-tester>

4 Clojure Solves ASDL Issues

4.1 Clojure Is a Complete Fit for ASR

ASR expressions, being trees, have a natural representation in S-Expressions.¹⁸ Clojure, being a modern Lisp, natively handles S-Expressions.

4.2 Clojure is not Moribund

Clojure is up-to-date, lively, and production-ready.

4.3 Clojure Has Type-Checking Tools

Clojure.spec,¹⁹ is a *force majeure* for precision, completeness, verification, and validation. The collection of MASR specs amounts to a meta-type system for ASR.

Clojure specs are arbitrary predicate functions. Clojure specs can easily express the difference between *list* and *set*, solving the ambiguity issue outlined in Section 3.5. Clojure specs, moreover, can flexibly express type-system features beyond the logics of typical, hard-coded type systems. That flexibility affords new long-term opportunities, say for experiments in dependent types and concurrency types.²⁰ In the short run, clojure.spec will make type constraints for ASDL explicit and manifest, and will relieve the burden on C++ programmers to manage implicit constraints.

4.4 ASR in Clojure Solves Volatility

We aim to replace ASDL with MASR in Clojure. When integrated with LCompilers in the future, MASR will be called LASR. There will be no gap or lag between LASR specs and their implementations because the implementations will be generated at build time.

¹⁸<https://en.wikipedia.org/wiki/S-expression>

¹⁹<https://clojuredocs.org/clojure.spec.alpha>

²⁰<https://rholang.io/>

5 MASR Definitions

Definition 1. A *spec* is a predicate function that tests an expression for conformance. *Spec* is a synonym for *type* in this document.

Definition 2. *Terms* are the "objects" or "productions" of ASR, like `symbol` or `dimension`.

Names of terms appear to the left of equals signs in the ASDL grammar snapshot.⁹ Names of terms are generally in lower-case.

Table 6 exhibits terms, ambiguous types, and term-like types, which are used-but-not-defined in the ASDL grammar. MASR explicitly defines them. Each term has a Clojure spec, various sugar functions, and term-entity key specs (Section 2.9.1).

The ambiguous types, `symbol_table` and `syntab_id`, are called out. The ASDL grammar conflates these two, having only `symbol_table`, which can either a full hash-map entity or an integer ID, depending on obscure criteria hidden in hand-written C++ code. A primary objective of MASR is to remove this kind of ambiguity, which is a design flaw in the current ASDL grammar rather than a deficiency of ASDL because ASDL can express the difference between a hash-map and an integer ID.

The contents of Table 6 have been greatly abbreviated and edited for presentation.

Definition 3. *Heads* are expressions like `Local` and `CaseStmt`, generally in `PascalCase`, that appear on the right-hand sides of equals signs in Table 6.

See the blog post in the footnote²¹ for an informal description of *PascalCase*.

There are two kinds of heads:

function-like heads — have parentheses and typed parameters,
e.g., `CaseStmt (expr*, stmt*)`

enum-like heads — no parentheses, e.g., `Local`

MASR has a Clojure spec and sugar functions for each head. Most heads also have term-head entity-key specs (Section 2.9.3). There are about 250 heads by a recent count.

Definition 4. An *ASR entity* is a compound type like `CaseStmt (expr*, stmt*)`, with a function-like head and zero-or more arguments, possibly with names, that require recursive conformance.

²¹<https://alok-verma6597.medium.com/case-styles-in-development-camel-pascal-snake-and-kebab-case-ed>

Table 6: Nodes in the ASDL Grammar

term	partial expansion
1 unit	TranslationUnit(symbol_table, node*)
2 symbol	...many heads...
3 storage_type	Default Save Parameter Allocatable
4 access	Public Private
5 intent	Local In Out InOut ...
6 deftype	Implementation Interface
7 presence	Required Optional
8 abi	Source LFortranModule ... Intrinsic
9 stmt	...many heads...
10 expr	...many heads...
11 ttype	Integer(int, dimension*) ...
12 restriction_arg	RestrictionArg(ident , symbol)
13 binop	Add Sub ... BitRShift
14 logicalbinop	And Or Xor NEqv Eqv
15 cmpop	Eq NotEq Lt LtE Gt GtE
16 integerboz	Binary Hex Octal
17 arraybound	LBound UBound
18 arraystorage	RowMajor ColMajor
19 cast_kind	RealToInteger IntegerToReal ...
20 dimension	(expr? start, expr? length)
21 alloc_arg	(expr a, dimension* dims)
22 attribute	Attribute(ident name, attr-arg* args)
23 attribute_arg	(ident arg)
24 call_arg	(expr? value)
25 tbind	Bind(string lang, string name)
26 array_index	(expr? left, expr? right, expr? step)
27 do_loop_head	(expr? v, expr? start expr? end, expr? step)
28 case_stmt	CaseStmt(expr*, stmt*) ...
29 type_stmt	TypeStmtName(symbol, stmt*) ...
30 enumtype	IntegerConsecutiveFromZero ...
ambiguous	
31 symbol_table	Clojure maps
32 symtab_id	int (new in MASR; not in ASDL)
*term-like	
0 dimensions	dimension*, via Clojure vectors or lists
0 atoms	int float bool nat bignat
0 identifier	by regex
0 identifiers	identifier*, via Clojure sets

6 MASR Tenets

Entity Hash-Maps — ASR entities¹² in full-form shall be hash-maps with fully-qualified keywords as keys (see Section 2.4 for a summary and Section 9.1 for a motivating example, *intent*).

Full-Form — Every entity shall have a fully explicit form in which all attributes are spelled out. Full-form hash-maps shall contain all necessary information, even at the cost of verbosity.

Multi-Specs — ASR entity hash-maps shall be recursively checked and generated via Clojure multi-specs. See Section 2.7 for discussion and the body of this document for examples.

Syntax Sugar — Certain constructor functions may supply default entity-attribute values so as to shorten common-case expressions. See Section 2.9 for discussion, Section 9.11 for an example, and see Issue 3 on MASR’s GitHub repo.²²

²²<https://github.com/rebcabin/masr/issues/3>

7 Base Specs

The specs in this section are the *atoms* in Table 1 and atoms in the *term-like* grouping in Table 6.

7.1 Atoms: `int`, `float`, `bool`, `nat`

The specs for `int`, `float`, and `bool` are straightforward:

```
(s/def ::int int?) ;; java.lang.Long
(s/def ::float float?)
(s/def ::bool boolean?)
```

7.1.1 Sugar

We restrict the spec, `nat`, for natural numbers, to `int`, for practical reasons:

```
(s/def ::nat nat-int?)
;; sugar
(defn nat [it]
  (let [cit (s/conform ::nat it)]
    (if (s/invalid? cit)
      ::invalid-nat
      cit)))
```

```
(tests
 (s/valid? ::nat (nat 42)) := true
 (s/valid? ::nat (nat -42)) := false
 (s/valid? ::nat (nat 0)) := true
 (s/valid? ::nat (nat 0xFFFFFFFFFFFFFFFF)) := false
 (s/valid? ::nat (nat -0xFFFFFFFFFFFFFFFF)) := false
 (s/valid?
  ::nat
  (nat (unchecked-long 0xFFFFFFFFFFFFFFFF))) := false
 (s/valid?
  ::nat
  (nat (unchecked-long -0xFFFFFFFFFFFFFFFF))) := true
 (s/valid? ::nat (nat 0x7FFFFFFFFFFFFFFFFF)) := true)
```

7.2 Notes

A Clojure *int* is a Java *Long*, with some peculiar behavior for hex literals.²³ Hex literals for negative numbers in Clojure must have explicit minus signs, lest they become `clojure.lang.BigInt`. MASR disallows `BigInt` for dimension (9.20) and dimensions (8.1). To get negative `java.lang.Long` without explicit minus signs, one employs Clojure's `unchecked-long`.

```
(tests (unchecked-long 0x8000000000000000)
      := -9223372036854775808
      (unchecked-long 0xFFFFFFFFFFFFFFFF)
      := -1
      (unchecked-long 0x8000000000000000)
      := -0x8000000000000000
      (unchecked-long -0xFFFFFFFFFFFFFFFF)
      := 1)
```

²³<https://clojurians.slack.com/archives/C03S1KBA2/p1681690965585429>

8 Term-Like Nodes

This section of the document exhibits specs for the *term-like nodes* in Tables 2 and 6: namely `dimensions` (plural), `identifier`, and `identifiers`. These are not terms, but share some similarities with terms. Note carefully the singulars and plurals in the names of the specs. `dimension` (singular) is a term and covered in Section 9.20, whereas `dimensions` (plural) is not a term. In the ASDL snapshot grammar,⁹ the plural of `dimension` is denoted `dimension*`, with a Kleene star.

8.1 dimensions [plural]

A MASR *dimensions* [plural], `dimension*` in ASDL, is a homogeneous ordered collection (list or vector) of zero or more `dimension` instances (9.20). Because `::dimensions` [plural] is not a term, we do not need nested multi-specs. However, because `::dimension` [singular] is a term, the elements of a `dimensions*` must conform to `::dimension`, which is an `asr-term` multi-spec. We ensure such conformance with a general-purpose function that selects terms that match a given spec, `dimension` in this case. MASR reuses that function in other specs that represent non-term collections.

```
(defn term-selector-spec [kwd]
  (s/and ::asr-term
    #(= kwd (::term %))))
```

Remark 1. The notation `#(...%...)` is Clojure shorthand for an anonymous function (lambda) with a positional argument denoted by `%`, and positional arguments `%1`, `%2`, ... when there are two or more arguments. Applying a keyword like `::term` as a function picks that keyword out of its hash-map argument.

Here is the spec, `::dimensions`, for dimensions. We limit the number of dimensions to 9 for practical reasons. The meaning of an empty `::dimensions` instance is an open question (Issue 7²⁴).

```
(def MIN-NUMBER-OF-DIMENSIONS 0) ;; TODO: 1?
(def MAX-NUMBER-OF-DIMENSIONS 9)

(s/def ::dimensions
  (s/coll-of (term-selector-spec ::dimension)
    :min-count MIN-NUMBER-OF-DIMENSIONS,
    :max-count MAX-NUMBER-OF-DIMENSIONS,
    :into []))
```

²⁴<https://github.com/rebcabin/masr/issues/7>

8.1.1 FullForm

The following tests show a couple of ways of writing out a `::dimensions` instance in full-form. The first is necessary in files other than `specs.clj`, say in `core_tests.clj`. The second can be used in `specs.clj`:

```
(tests (s/valid?
  ::dimensions
  [#:masr.specs{:term :masr.specs/dimension,
    :dimension-content [1 60]}
   #:masr.specs{:term :masr.specs/dimension,
    :dimension-content ()}]) := true
(s/valid?
  ::dimensions
  [{::term ::dimension,
    :dimension-content [1 60]}
   {::term ::dimension,
    :dimension-content ()}]) := true)
```

8.1.2 Sugar

The following tests illustrate the sugar function, `dimensions`, for the spec, `::dimensions`:

```
(tests
  (s/valid? ::dimensions []) := true
  (s/valid? ::dimensions
    [(dimension '(1 60)) (dimension '())]) := true
  (s/conform ::dimensions
    [(dimension '(1 60)) (dimension '())]) :=
  [#:masr.specs{:term :masr.specs/dimension,
    :dimension-content [1 60]}
   #:masr.specs{:term :masr.specs/dimension,
    :dimension-content ()}])
```

8.2 identifier [*singular*]

An ASR identifier is a C or Fortran identifier, which begins with an alphabetic glyph or an underscore, and has alpha-numeric characters or underscores following. The only complication in the spec is the need to generate instances. The spec solves the generation problem for identifiers, plus shows a pattern for other specs that need custom generators.

```
(let [alpha-re #"[a-zA-Z_]" ;; "let over lambda."
      alphameric-re #"[a-zA-Z0-9_]*"]
  (def alpha?
    #(re-matches alpha-re %))
  (def alphameric?
    #(re-matches alphameric-re %))
  (defn identifier? [sy]
    ;; exclude strings, numbers, quoted numbers
    (and (symbol? sy)
          (let [s (str sy)]
            (and (alpha? (subs s 0 1))
                  (alphameric? (subs s 1))))))
  (def identifier-generator
    (tgen/let [c (gen/char-alpha)
              s (gen/string-alphanumeric)]
      (symbol (str c s))))
  (s/def ::identifier
    (s/with-gen
      identifier?
      ;; fn wrapping a macro:
      (fn [] identifier-generator))))
```

The following tests illustrate validation and generation:

```
(tests
  (s/valid? :masr.specs/identifier 'foobar) := true
  (s/valid? :masr.specs/identifier '_f__547) := true
  (s/valid? :masr.specs/identifier '1234)    := false)
#_
(gen/sample (s/gen :masr.specs/identifier))
;; => (e c Q G Z2qP fXzg1 sRx2J6 YIhKlV k6 f7k1Xl4)
;; => (k hM LV QWC qW0X RGk3u W Kg6X Q2YvFO621 ODUt9)
```

8.2.1 Sugar

We define and illustrate the sugar function, `identifier` for creating identifiers:

```
(defn identifier [sym]
  (let [csym (s/conform ::identifier sym)]
    (if (s/invalid? csym)
        ::invalid-identifier
        csym)))
(tests
 (identifier 'foo)  := 'foo
 (identifier 123)   := ::invalid-identifier)
```

8.3 identifiers [*plural*]

ASDL `identifier*` is ambiguous. There are three kinds of identifier collections in MASR:²⁵

identifier-set — unordered, no duplicates

identifier-list — ordered, duplicates allowed (we use vector)

identifier-suit — ordered, duplicates not allowed

For all three kinds, we limit the number of identifiers to 99 for practical purposes:

```
(def MIN-NUMBER-OF-IDENTIFIERS 0)
(def MAX-NUMBER-OF-IDENTIFIERS 99)
```

8.3.1 identifier-set

The spec for a set of identifiers is straightforward because of Clojure's literal syntax, `#{\ldots}`, for sets, including the empty set:

```
(s/def ::identifier-set
  (s/coll-of ::identifier
    :min-count MIN-NUMBER-OF-IDENTIFIERS,
    :max-count MAX-NUMBER-OF-IDENTIFIERS,
    :into #{})) ;; empty set
```

See the code for uninteresting details of the sugar function, `identifier-set`. The following tests show it at work:

```
(tests
  (let [x (identifier-set ['a 'a])]
    (s/valid? ::identifier-set x) := true
    (set? x) := true
    (count x) := 1)
  (let [x (identifier-set [])]
    (s/valid? ::identifier-set x) := true
    (set? x) := true
    (count x) := 0)
  (let [x (identifier-set ['a '1])]
    (s/valid? ::identifier-set x) := false
    x := ::invalid-identifier-set))
```

²⁵<https://github.com/rebcabin/masr/issues/1>

8.3.2 identifier-list

The spec for a list of identifiers is almost the same as the spec for a set of identifiers. It differs only in the `:into` clause — into a vector rather than into a set:

```
(s/def ::identifier-list
  (s/coll-of ::identifier
    :min-count MIN-NUMBER-OF-IDENTIFIERS,
    :max-count MAX-NUMBER-OF-IDENTIFIERS,
    :into []))

(tests
  (every? vector? (gen/sample
    (s/gen ::identifier-list))) := true)
```

The implementation of the sugar function for `identifier-list` is uninteresting. The following tests show it at work:

```
(tests
  (let [x (identifier-list ['a 'a])]
    (s/valid? ::identifier-list x) := true
    (vector? x) := true
    (count x) := 2)
  (let [x (identifier-list [])]
    (s/valid? ::identifier-list x) := true
    (vector? x) := true
    (count x) := 0)
  (let [x (identifier-list ['a '1])]
    (s/valid? ::identifier-list x) := false
    x := ::invalid-identifier-list))
```

8.3.3 identifier-suit

The spec for an identifier-suit is almost the same as for identifier-list, only checking that there are no duplicate elements:

```
(s/def ::identifier-suit
  (s/and
    (s/coll-of ::identifier
      :min-count MIN-NUMBER-OF-IDENTIFIERS,
      :max-count MAX-NUMBER-OF-IDENTIFIERS,
      :into [])
    ;; no duplicates
    #(= (count %) (count (set %)))))
```

Here are the tests for the (uninteresting) sugar function:

```
(tests
  (let [x (identifier-suit ['a 'a])]
    (s/valid? ::identifier-suit x) := false
    (vector? x) := false)
  (let [x (identifier-suit ['a 'b])]
    (s/valid? ::identifier-suit x) := true
    (vector? x) := true
    (count x) := 2)
  (let [x (identifier-suit [])]
    (s/valid? ::identifier-suit x) := true
    (vector? x) := true
    (count x) := 0)
  (let [x (identifier-suit ['a 'l])]
    (s/valid? ::identifier-suit x) := false
    x := ::invalid-identifier-suit))
```

9 Specs

The following sections

- summarize the Clojure specs for all ASR terms and heads (see Tables 4 and 5).
- pedagogically explain the architecture and approach taken in the Clojure code so that anyone may extend and maintain it.

The architecture is the remainder from several experiments. For example, `defrecord` and `defprotocol` for polymorphism were tried and discarded in favor of multi-specs.

The tests in `core_test.clj` exhibit many examples that pass and, more importantly, fail the specs. We also keep lightweight, load-time tests inline to the source file, `specs.clj`. We don't have strict criteria for whether a test should be inline, separate, or both.

The best way to learn the code is to study the tests and to run them in the Clojure REPL or in the CIDER debugger in Emacs.²⁶

We present the terms somewhat out of the order of Table 6. First is *intent*, as it is the archetype for several enum-like terms and heads.

²⁶<https://docs.cider.mx/cider/debugging/debugger.html>

9.1 intent

9.1.1 Sets for Contents

An ASR *intent* is one of the symbols

Local, In, Out, InOut, ReturnVar, Unspecified.

The spec for the *contents* of an intent is simply this set of enum-like heads. Any Clojure *set* (e.g., in `#{ ... }` brackets) doubles as a predicate function for set membership. In the following two examples, the set appears in the function position of the usual Clojure function-call syntax (*function args**):

If a candidate member is in a set, the result of calling the set like a function is the candidate member:

```
(#{'Local 'In 'Out 'InOut 'ReturnVar 'Unspecified} 'Local)
```

Local

When the candidate element, say *fubar*, is not in the set, the result is *nil*, which does not print:

```
(#{'Local 'In 'Out 'InOut 'ReturnVar 'Unspecified} 'fubar)
```

Any predicate function can be registered as a Clojure spec.¹⁹ Therefore the spec for *intent contents* is just the set of valid members.

9.1.2 Specs have Fully Qualified Keyword Names

The name of the spec is `::intent-enum`. The double colon in `::intent-enum` is shorthand. In the file `specs.clj`, double colon implicitly signifies that a keyword like `::intent-enum` is in the namespace `masr.specs`. In other files, like `core_test.clj`, the same keyword is spelled `:masr.specs/intent-enum`.

The names of all Clojure specs must be fully qualified in namespaces.

```
(s/def ::intent-enum
  #{'Local 'In 'Out 'InOut 'ReturnVar 'Unspecified})
```

9.1.3 How to Use Specs

To check an expression like `'Local` against the `::intent-enum` spec, write

```
(s/valid? ::intent-enum 'Local)
;; => true
(s/valid? ::intent-enum 'fubar)
;; => false
```

To produce conforming or non-conforming (invalid) entities in other code, write

```
(s/conform ::intent-enum 'Local)
;; => Local
(s/conform ::intent-enum 'fubar)
;; => :clojure.spec.alpha/invalid
```

To generate a few conforming samples, write

```
(gen/sample (s/gen ::intent-enum) 5)
;; => (Unspecified Unspecified Out Unspecified Local)
```

or, with conformance explanation (trivial in this case):

```
(s/exercise ::intent-enum 5)
;; => ([Out Out]
;;      [ReturnVar ReturnVar]
;;      [In In]
;;      [Local Local]
;;      [ReturnVar ReturnVar])
```

Strip out the conformance information as follows:

```
(map second (s/exercise ::intent-enum 5))
;; => (In ReturnVar Out In ReturnVar)
```

`s/valid?`, `s/conform`, `gen/sample`, and `s/exercise` pertain to any Clojure specs, no matter how complex or rich.

9.1.4 The Spec that Contains the Contents

`::intent-enum` is just the spec for the *contents* of an intent, not for the intent itself. The spec for the intent itself is an implementation of a polymorphic Clojure *multi-spec*,¹² `::asr-term`.

9.1.5 Multi-Specs

A multi-spec is like a tagged union in C. The multi-spec, `::asr-term`, pertains to all Clojure hash-maps¹⁵ that have a tag named `::term` with a value like `::intent` or `::storage-type`, etc. The values, if themselves fully qualified keywords, are recursively checked.

A multi-spec has three components:

defmulti¹⁶ — a polymorphic interface that declares the *tag-fetcher function*, `::term` in this case. The tag-fetcher function fetches a tag's value from any candidate hash-map. The `defmulti` dispatches to a `defmethod` that matches the fetched tag value, `::intent` in this case. `::term` is a fully qualified keyword of course, but all keywords double as tag-fetchers for hash-maps.²⁷

defmethod²⁸ — individual specs, each implementing the interface. In this case, if the `::term` of a hash-map matches `::intent`, then the corresponding `defmethod` is invoked (see Section 9.1.7 below).

s/multi-spec — tying together the `defmulti` and, redundantly, the tag-fetcher.²⁹

9.1.6 Specs for All Terms

Start with a spec for `::term`:

```
;; like ::intent, ::symbol, ::expr, ...  
(s/def ::term qualified-keyword?)
```

The spec says that any fully qualified keyword, like `::intent`, is a MASR term. This spec leaves room for growth of MASR by adding more fully qualified keywords for more MASR types-*qua*-terms.

`s/def` stands for `clojure.spec.alpha/def`, the `def` macro in the `clojure.spec.alpha` namespace. The namespace is aliased to `s`.

Next, specify the `defmulti` polymorphic interface, `term`, (no colons) for all term specs, and bind it to the tag-fetcher, `::term`:

```
(defmulti term ::term)
```

This `defmulti` dispatches to a `defmethod` based on the results of applying the keyword-*qua*-tag-fetcher function `::term` to a hash-map:

```
(::term {::term ::intent ...})  
;; => ::intent
```

The spec is named `::term` and the tag-fetcher is named `::term`. They don't need to be the same name, but they always are in MASR.

²⁷<https://stackoverflow.com/questions/6915531>

²⁸<https://clojuredocs.org/clojure.core/defmethod>

²⁹Multi-specs allow re-tagging, where the tag named in the multi-spec is different from the tag-fetcher function. MASR does not need re-tagging.

9.1.7 Spec for intent

If applying `::term` to a Clojure hash-map produces `::intent`, the following spec, which specifies all intents, will be invoked. It ignores its argument, `_`:

```
(defmethod term ::intent [_]
  (s/keys :req [::term ::intent-enum]))
```

This spec states that an *intent* is a Clojure hash-map with a required `::term` keyword and a required `::intent-enum` keyword. This kind of spec is called an *entity spec*.¹²

9.1.8 The Multi-Spec Itself: `::asr-term`

`s/multi-spec` ties `defmulti` `term` to the tag-fetcher `::term`. The multi-spec itself is named `::asr-term`:

```
;;      name of the mult-spec      defmulti  tag fn
;;      -----
(s/def ::asr-term (s/multi-spec  term      ::term))
```

9.1.9 Examples of Intent

The following shows a valid example:

```
(s/valid? ::asr-term
  {::term      ::intent,
   ::intent-enum 'Local})
```

true

Here is an invalid sample:

```
(s/valid? ::asr-term
  {::term      ::intent,
   ::intent-enum 'FooBar})
```

false

Generate a few valid samples:

```
(gen/sample (s/gen (s/and
  ::asr/asr-term
  #(= ::asr/intent (::asr/term %))))
  5)
;;=> (::asr{:term ::asr/intent, :intent-enum ReturnVar}
;;    (::asr{:term ::asr/intent, :intent-enum In}
;;    (::asr{:term ::asr/intent, :intent-enum Unspecified}
;;    (::asr{:term ::asr/intent, :intent-enum Unspecified}
;;    (::asr{:term ::asr/intent, :intent-enum InOut})))
```

9.1.10 Another asr-term: a Pattern Emerges

To define another asr-term, specify the contents and write a defmethod. The one multi-spec, `::asr-term`, suffices for all.

For example, another asr-term for an enum-like is `storage-type`:

```
(s/def ::storage-type-enum
  #{'Default, 'Save, 'Parameter, 'Allocatable})

(defmethod term ::storage-type [_]
  (s/keys :req [::term ::storage-type-enum]))
```

All enum-like specs follow this pattern.

9.1.11 Syntax Sugar

`{::term ::intent, ::intent-enum 'Local}`, a valid asr-term entity, is long and ugly. Write a short function, `intent`, via `s/conform`, explained in Section 9.1.3:

```
(defn intent [sym]
  (let [intent_ (s/conform
    ::asr-term
    {::term ::intent, ::intent-enum sym})]
    (if (s/invalid? intent_)
      ::invalid-intent
      intent_)))
```


Entities have shorter expression with the sugar:

```
(testing "better syntax"
  (is (s/valid? ::asr-term (intent 'Local)))
  (is (s/valid? ::asr-term (intent 'Unspecified)))
  (is (not (s/valid? ::asr-term (intent 'foobar))))
  (is (not (s/valid? ::asr-term (intent []))))
  (is (not (s/valid? ::asr-term (intent ())))
  (is (not (s/valid? ::asr-term (intent {}))))
  (is (not (s/valid? ::asr-term (intent #{ }))))
  (is (not (s/valid? ::asr-term (intent "foobar"))))
  (is (not (s/valid? ::asr-term (intent ""))))
  (is (not (s/valid? ::asr-term (intent 42))))
  (is (thrown? clojure.lang.ArityException (intent))))
```

All our specs are like that: a long-form hash-map and a short-form sugar function that does a conformance check.

9.1.12 Capture the Enum-Like Pattern in a Macro

All enum-likes have a *contents* spec, a `defmethod term`, and a syntax-sugar function. The following macro pertains to all such enum-like multi-specs:

```
(defmacro enum-like [term, heads]
  (let [ns "masr.specs"
        tkw (keyword ns (str term))
        tke (keyword ns (str term "-enum"))
        tki (keyword ns (str "invalid-" term))]
    `(do
      (s/def ~tke ~heads) ;; the set
      (defmethod term ~tkw [_#] ;; the multi-spec
        (s/keys :req [:masr.specs/term ~tke]))
      (defn ~term [it#] ;; the sugar
        (let [st# (s/conform
                    :masr.specs/asr-term
                    {:masr.specs/term ~tkw
                     ~tke it#})]
          (if (s/invalid? st#) ~tki, st#))))))
```

Use the macro like this:

```
(enum-like
 intent
 #{'Local 'In 'Out 'InOut 'ReturnVar 'Unspecified})
(enum-like
 storage-type
 #{'Default, 'Save, 'Parameter, 'Allocatable})
```

9.1.13 Term-Entity Keys

The actual `enum-like` macro also defines the *term-entity-key spec* (Section 2.9.1) for any enum-like.

```
(s/def ~tkw    ;; like ::intent
  (s/and ~art ;; like ::asr-term, i.e., the multi-spec
    ;; like the predicate #(= ::intent (::term %))
    (term-selector-spec ~tkw)))
```

In this case, the term-entity-key spec is `::intent`:

```
(testing "term entity-key"
  (is (s/valid? ::asr/intent (intent 'Local)))
  (is (s/valid? ::asr/intent (intent 'Unspecified)))
  (is (not (s/valid? ::asr/intent (intent 'foobar))))
  (is (not (s/valid? ::asr/intent (intent []))))
  (is (not (s/valid? ::asr/intent (intent ())))))
  (is (not (s/valid? ::asr/intent (intent {}))))
  (is (not (s/valid? ::asr/intent (intent #{}))))
  (is (not (s/valid? ::asr/intent (intent "foobar"))))
  (is (not (s/valid? ::asr/intent (intent ""))))
  (is (not (s/valid? ::asr/intent (intent 42))))
  (is (thrown? clojure.lang.ArityException
    (intent))))
```

9.2 **TODO** unit

9.3 TODO symbol

9.3.1 Variable

Here is an example of the full form for a `::Variable` with a conforming instance in light sugar (Section 2.9). Note the term-head-entity-key spec `::Variable` (Section 2.9.3). Any `::Variable` is also an `::asr-term`.

```
(let [a-var-head
      {::symbol-head ::Variable
       ::syntab-id    (nat 2)
       ::varnym       (varnym 'x)
       ::ttype        (ttype (Integer 4 []))
       ::type-declaration (type-declaration nil)
       ::dependencies  (identifier-set ())
       ::intent        (intent 'Local)
       ::symbolic-value () ;; TODO sugar
       ::value         () ;; TODO sugar
       ::storage-type  (storage-type 'Default)
       ::abi            (abi 'Source :external false)
       ::access         (access 'Public)
       ::presence       (presence 'Required)
       ::value-attr     false ;; TODO sugar
      }
      a-var {::term ::symbol
             ::asr-symbol-head a-var-head}
      a-var-light (Variable-
                   :varnym      (identifier 'x)
                   :syntab-id   2
                   :ttype       (ttype (Integer 4))))]
```

```
(tests
 a-var-light := (s/conform ::asr-term a-var)
 a-var-light := (s/conform ::Variable a-var)

 (s/valid? ::asr-symbol-head a-var-head) := true

 (s/valid? ::asr-term a-var)              := true
 (s/valid? ::asr-term a-var-light)        := true

 (s/valid? ::Variable a-var)              := true
 (s/valid? ::Variable a-var-light)        := true))
```

Here is an example in heavy sugar:

```
(let [a-valid (Variable 2 'x (Integer 4)
                           nil [] 'Local
                           [] [] 'Default
                           'Source 'Public 'Required
                           false)]
  (s/valid? ::asr-term a-valid) := true
  (s/valid? ::Variable a-valid) := true)
```

The source file, `specs.clj`, tests each of the 13 positional arguments of the heavy-sugar function `Variable` for recursive conformance.

9.4 storage_type

Storage-type is another enum-like, defined and registered via macro (Section 9.1.12). The following tests of full-form and heavy sugar illustrate conformance to both `::asr-term` and the term-entity-key spec, `::storage-type` (Section 2.9.1).

```
(tests
  (s/valid? ::storage-type-enum 'Default)           := true
  (s/valid? ::storage-type-enum 'foobar)            := false
  (s/valid? ::asr-term
    {::term ::storage-type
      ::storage-type-enum 'Default})                := true
  (s/valid? ::asr-term (storage-type 'Default))      := true
  (s/valid? ::asr-term (storage-type 'foobar))       := false
  (s/valid? ::storage-type
    {::term ::storage-type
      ::storage-type-enum 'Default})                := true
  (s/valid? ::storage-type (storage-type 'Default)) := true
  (s/valid? ::storage-type (storage-type 'foobar)) := false
  (storage-type 'foobar)                             := ::invalid-storage-type
  (let [ste (storage-type 'Default)]
    (s/conform ::storage-type ste)                    := ste
    (s/conform ::asr-term ste)                        := ste))
```

9.5 access

Access is another enum-like, defined and registered via macro (Section 9.1.12). The following tests of heavy sugar illustrate conformance to both `::asr-term` and the term-entity-key spec, `::access` (Section 2.9.1).

```
(enum-like access #{'Public 'Private})

(tests
  (let [public (access 'Public)]
    (s/conform ::asr-term public) := public
    (s/conform ::access public) := public)
  (access 'foobar) := ::invalid-access)
```

9.6 TODO deftype

9.7 presence

Presence is another enum-like, defined and registered via macro (Section 9.1.12). The following tests of heavy sugar illustrate conformance to both `::asr-term` and the term-entity-key spec, `::presence` (Section 2.9.1).

```
(enum-like presence #{'Required 'Optional})

(tests
  (let [required (presence 'Required)]
    (s/conform ::asr-term required) := required
    (s/conform ::presence required) := required)
  (presence 'fubar) := ::invalid-presence)
```

9.8 abi

Abi is a rich case. It is enum-like, similar to *intent* (Section 9.1), but with restrictions. Its heads include several *external-abis*:

```
(def external-abis
  #{ 'LFortranModule, 'GFortranModule,
    'BindC, 'Interactive, 'Intrinsic})
```

and one *internal-abi*:

```
(def internal-abi #{ 'Source})
```

The *abi-enum* contents spec for *abi* is the union of these two sets:

```
(s/def ::abi-enum
  (set/union external-abis internal-abis))
```

Specify an additional key in a conforming *abi* hash-map with a `::bool` predicate:

```
(s/def ::abi-external ::bool)
```

Add a convenience function for logic:

```
(defn iff [a b]
  (or (and a b)
      (not (or a b))))
```

Specify the `defmethod` for the *abi* itself with a hand-written generator (clojure.spec is not strong enough to create the generator automatically):

```
(defmethod term ::abi [_]
  (s/with-gen
    (s/and
      # (iff (= 'Source (::abi-enum %))
              (not (::abi-external %)))
      (s/keys :req [::term ::abi-enum ::abi-external]))
    (fn []
      (tgen/one-of
        [ (tgen/hash-map
            ::term      (gen/return ::abi)
            ::abi-enum  (s/gen external-abis)
            ::abi-external (gen/return true))
          (tgen/hash-map
            ::term      (gen/return ::abi)
            ::abi-enum  (s/gen internal-abis)
            ::abi-external (gen/return false))] ))))
```


Generate a few conforming samples in full-form:

```
(gen/sample (s/gen (s/and
                  ::asr/asr-term
                  #(= ::asr/abi (::asr/term %))))
            5)
;; => (::asr{:term ::asr/abi,
;;      :abi-enum Interactive, :abi-external true}
;;      (::asr{:term ::asr/abi,
;;      :abi-enum Source, :abi-external false}
;;      (::asr{:term ::asr/abi,
;;      :abi-enum Source, :abi-external false}
;;      (::asr{:term ::asr/abi,
;;      :abi-enum Source, :abi-external false}
;;      (::asr{:term ::asr/abi,
;;      :abi-enum Interactive, :abi-external true}))
```

The sugar for *abi* is an exceptional case. We deem it better to default the `:external` Boolean to `false` in all cases except that for `'Source`, and to require an explicit `:external` keyword in other cases. That means that arity-1 usages like

```
(abi 'Source)
```

and arity-3 usages like

```
(abi 'Source :external false)
```

are valid, but arity-2 usages like

```
(abi 'Source false)
```

are not valid.

The following sugar function effects this design:

```
(defn abi
  ;; arity 1 --- default "external"
  ([the-enum]
   (let [abi_ (s/conform
                 ::asr-term
                 {::term      ::abi,
                  ::abi-enum  the-enum,
                  ::abi-external
                    (not (= the-enum 'Source))})]
     (if (s/invalid? abi_)
         ::invalid-abi
         abi_)))
  ;; arity 2 --- invalid
  ([the-enum, crap]
   ::invalid-abi)
  ;; arity 3 --- light sugar
  ([the-enum, ext-kw, the-bool]
   (cond
     (not (= ext-kw :external)) ::invalid-abi
     :else
     (let [abi_ (s/conform
                   ::asr-term
                   {::term      ::abi,
                    ::abi-enum  the-enum,
                    ::abi-external the-bool})]
       (if (s/invalid? abi_)
           ::invalid-abi
           abi_))))))
```

Here is its term-entity key, `::abi`, for recursive checking (Section 2.9.1):

```
(def-term-entity-key abi)
```

Here are some conformance tests for full-form, sugar, against `::asr-term`, and against the term-entity key `::abi`:

```
(tests
 (s/valid? ::asr-term
  {::term      ::abi
   ::abi-enum 'Source
   ::abi-external false})      := true
 (let [abe (abi 'Source :external false)]
  (s/conform ::abi      abe)      := abe
  (s/conform ::asr-term abe)      := abe
  ;; defaults to correct value
  (abi 'Source)                  := abe
  ;; missing keyword
  (abi 'Source false)            := ::invalid-abi
  ;; wrong value
  (abi 'Source :external true)   := ::invalid-abi)
 (let [abe (abi 'LFortranModule :external true)]
  (s/conform ::asr-term abe)      := abe
  (s/conform ::abi      abe)      := abe
  ;; defaults to correct value
  (abi 'LFortranModule)          := abe
  ;; missing keyword
  (abi 'LFortranModule true)     := ::invalid-abi
  ;; wrong value
  (abi 'LFortranModule :external false) := ::invalid-abi))
```

9.9 **TODO** stmt

9.10 **TODO** expr

9.11 ttype

Ttype [sic] features a nested multi-spec. Ttype is an archetype for all function-like heads, just as *intent* is an archetype for all enum-like heads.

```
(defmulti ttype-head ::ttype-head)
(defmethod ttype-head ::Integer [_]
  (s/keys :req [::ttype-head ::bytes-kind ::dimensions]))
(s/def ::asr-ttype-head
  (s/multi-spec ttype-head ::ttype-head))
```

```
(defmethod term ::ttype [_]
  (s/keys :req [::term ::asr-ttype-head]))
```

9.11.1 Full Form

One may always write out ttype specs in full:

```
(s/valid? ::asr-term
  {::term ::ttype,
   ::asr-ttype-head
   {::ttype-head ::Integer,
    ::bytes-kind 4
    ::dimensions [[6 60] [1 82]]}})
```

9.11.2 Sugar for Integer, Real, Complex, Logical

Sugar for ttypes comes in two varieties, *light sugar* and *heavy sugar*. See Section 2.9 for rationale.

1. Light Sugar Examples

```
(ttype (Integer- {:dimensions [], :kind 4}))
(ttype (Integer- {:kind 4, :dimensions []}))
```

2. Heavy Sugar Examples

```
(ttype (Integer))
(ttype (Integer 4))
(ttype (Integer 2 []))
(ttype (Integer 8 [[6 60] [1 42]]))
```

See the tests for many more examples.

9.11.3 TODO Character

9.12 TODO restriction_arg

9.13 **TODO** binop

9.14 **TODO** logicalbinop

9.15 **TODO** cmpop

9.16 **TODO** integerboz

9.17 **TODO** arraybound

9.18 **TODO** arraystorage

9.19 **TODO** cast_kind

9.20 dimension

A *dimension* [*singular*] is 0 or 2 nats in a Clojure list or vector:

```
(def MIN-DIMENSION-COUNT 0)
(def MAX-DIMENSION-COUNT 2)
(s/def ::dimension-content
  (s/and (fn [it] (not (= 1 (count it))))
    (s/coll-of ::nat
      :min-count MIN-DIMENSION-COUNT,
      :max-count MAX-DIMENSION-COUNT,
      :into ())))
```

If there are two nats, the first nat specifies the starting index of any array dimension that enjoys the instance, and the second nat specifies the length. For example, in the ttype `(Integer 4 [[6 60]])` (9.11), the one dimension [*singular*] in the dimensions [*plural*] (8.1) of the ttype is `[6 60]`. The ttype specifies a rank-1 array of 60 4-byte integers with indices starting at 6 and running through 65.

If there are no nats, i.e., the array dimension of any array enjoying the instance is of zero length. For an example, consider the ttype `(Integer 4 [[]])` (9.11). This meaning of this type is an open question.²⁴

9.20.1 Empty Dimensions

Empty dimensions [*plural*], as in `(Integer 4 [])` specify non-array types. These are often called, loosely, *scalars*. Pedantically, *scalars* pertain only to a vector space.

An empty dimension, as in `(Integer 4 [[]])`, specifies a rank-1 array of zero length. Such items are discussed further in Issue 7²⁴ and in Section 8.1.

9.20.2 TODO: Issue 7: Zero Length

The following specs, in context of a ttype (9.11) for convenience, are legal in the ASDL grammar.⁹ They all denote arrays of length 0, and the meaning of an array of length 0 is **unspecified**:

```
(Integer 4 [[]])
(Integer 4 [[0]])
(Integer 4 [[6 0]])
```

9.20.3 FullForm

The following tests illustrate the full form for *dimension*:

```
(tests
  (s/valid? ::asr-term
    {::term ::dimension
      ::dimension-content [6 60]}) := true
  (s/valid? ::asr-term
    {::term ::dimension
      ::dimension-content [0]})   := false
  (s/valid? ::asr-term
    {::term ::dimension
      ::dimension-content []})    := true)
```

9.20.4 Sugar

The following tests illustrate the sugar and the term-entity-key spec (Section 2.9.1) for *dimension*:

```
(tests
  (s/conform ::asr-term
    {::term ::dimension,
      ::dimension-content '(1 60)}) :=
  (dimension '(1 60))
  (s/valid? ::asr-term (dimension 60))           := false
  (s/valid? ::asr-term (dimension [[]]))          := false
  (s/valid? ::asr-term (dimension 'foobar))        := false
  (s/valid? ::asr-term (dimension ['foobar']))     := false
  ;; throw arity (s/valid? ::asr-term (dimension)) := false
  (s/valid? ::asr-term (dimension []))             := true
  (s/valid? ::asr-term (dimension [60]))           := false
  (s/valid? ::asr-term (dimension [0]))            := false
  (s/valid? ::asr-term (dimension '(1 60)))        := true
  (s/valid? ::asr-term (dimension '()))            := true

  (s/valid? ::dimension (dimension 60))           := false
  (s/valid? ::dimension (dimension [[]]))          := false
  (s/valid? ::dimension (dimension 'foobar))        := false
  (s/valid? ::dimension (dimension ['foobar']))     := false
  (s/valid? ::dimension (dimension []))             := true
  (s/valid? ::dimension (dimension [60]))           := false
  (s/valid? ::dimension (dimension [0]))            := false
  (s/valid? ::dimension (dimension '(1 60)))        := true
  (s/valid? ::dimension (dimension '()))            := true )
```

9.21 **TODO** alloc_arg

9.22 **TODO** attribute

9.23 **TODO** attribute_arg

9.24 **TODO** call_arg

9.25 **TODO** tbind

9.26 **TODO** array_index

9.27 **TODO** do_loop_head

9.28 **TODO** case_stmt

9.29 **TODO** type_stmt

9.30 **TODO** enumtype

10 Implicit Terms

Terms used, explicitly or implicitly, but not defined in ASDL.

Some items specified in ASDL as *symbol_table* are actually *syntab_id*.

10.1 syntab_id

```
(s/def ::syntab-id ::nat)
```

```
(tests
  (syntab-id -42)           := ::invalid-syntab-id
  (syntab-id 'foo)         := ::invalid-syntab-id
  (syntab-id 42)           := 42
  (s/conform ::nat 42)      := 42
  (s/conform ::nat (nat 42)) := 42
  (s/conform ::syntab-id 42) := 42
  (s/conform ::syntab-id (syntab-id 42)) := 42
  (s/conform ::syntab-id (nat 42)) := 42)
```

10.2 TODO symbol_table

```
(s/def ::symbol-table map?
```